

Data-based, synthesis-driven: setting the agenda for computational ecology

Timothée Poisot ^{1,2}

¹: Université de Montréal, Département de Sciences Biologiques; ²: Québec Centre for Biodiversity Sciences

Writing manuscripts doesn't have to be awful. This template *tries* to provide good defaults for both the draft and preprint modes. Most of the information is contained in the YAML file. The preprint mode is particularly good looking.

Keywords: ecological networks - beta-diversity - biogeography



This work is licensed under a Creative Commons Attribution 4.0 Unported License.

Correspondence to Timothée Poisot – timothee.poisot@umontreal.ca — Latest update on May 5, 2017

Computational science happens when algorithms, software, data management practices, and advanced research computing are put in interaction with the explicit goal of solving “complex” problems. Typically, problems are considered *complex* when they cannot be solved appropriately with modelling or data-collection only. Computational science is one of the ways to practice computational thinking (Papert 1996), *i.e.* the feedback loop of abstracting a problem to its core mechanisms, expressing a solution in a way that can be automated, and using interactions between simulations and data to refine the original problem or suggest new knowledge. Computational approaches are common place in most area of biology, to the point where one would almost be confident that they represent a viable career path (Bourne 2011). Data usually collected in ecological studies have a high variability, are time-consuming, costly, and demanding to collect. In parallel, many problems lack appropriate formal mathematical formulations. For these reasons, computational approaches hold great possibilities, notably to further ecological synthesis and help decision-making (Petrovskii & Petrovskaya 2012).

Levin (2012) suggested that ecology (and evolutionary biology) should continue their move towards a *marriage of theory and data*. In addition to the aforementioned problem of the lack of adequately expressed models, this effort is hampered by the fact that data and models are often developed in independent ways, and reconciling both can be difficult. This has been suggested as one of the reasons for which theoretical papers (defined as *papers with at least one equation in the main text*) experience a sharp decrease in citation (Fawcett & Higginson 2012); this is the tragic sign that empirical scientists do not see the value of theoretical work, which of course can be blamed on both parties. One of the leading textbooks for the mathematical models in ecology and evolution (Otto & Day 2007) is more focused with algebra and calculus, and not with the integration of models with data. Other manuals that cover the integration of models and data tend to lean more towards statistical models (Bolker 2008; Soetaert & Herman 2008). This paints a picture of ecology as a field in which dynamical models and empirical data do not interact much, and instead the literature develops in silos.

Ecology as a whole (and community ecology in particular) circumvented the problem of model and data mismatch by investing in the development and refinement of statistical models (see Warton et al. 2014 for an excellent overview) and “numerical” approaches (Legendre & Legendre 1998) based on multivariate statistics. These models, however, are able to *explain* data, but very rarely do they give new predictions. This is, essentially, the niche that computational ecology can fill; at the cost of a higher degree of abstraction, its integration of data and generative models (*i.e.* models that, given rules, will generate new data) can be helpful to initiate the investigation of questions that have not received extensive empirical treatment, or for which usual statistical approaches fall short.

What is computational ecology? It is the application of computational thinking to ecological problems. This defines three core characteristics of computational ecology. First, it recognizes ecological systems as complex and adaptive; this places a great emphasis on mathematical tools that can handle, or even require, a certain degree of stochasticity (Zhang 2010, 2012). Second, it understands that data are the final arbiter of any simulation or model (Petrovskii & Petrovskaya 2012); this favours the use of data-driven approaches and analyses (Beaumont 2010). Finally, it accepts that some ecological systems are too complex to be formulated in mathematical or programmatic terms (Pascual 2005); the use of conceptual, or “toy” models, as long as they can be confronted to empirical data, is preferable to “abusing” mathematics by describing the wrong mechanism well (May 2004).

In a thought-provoking essay, Markowitz (2017) suggests that *all biology is computational biology* – the rationale behind this bold statement being that integrating computational advances, novel mathematical tools, and the usual data from one field, has a high potential to deliver synthesis. A more reasonable statement would be that *all ecology can benefit from computational ecology*, as long as we can understand how it interacts with other approaches; in this paper, I attempt to situate the practice of computational ecology within the broader landscape of ecological research. In particular, I highlight the ways in which computational ecology differs from, and complements, ecological modelling. I finally move on to the currency of collaborations between different groups of ecologists, and discuss the need to add more quantitative skills in ecological training.

1 A SUCCESS STORY: SPECIES DISTRIBUTION MODELS

The practice known as “species distributions modelling” (and the species distribution models, henceforth SDMs, it generates) is a good example of computational practices generating novel ecological insights. At their core, SDMs seek to model the presence (and sometimes absence) of a species based on previous observations of its presence (and sometimes absence) and knowledge of the environment in which the observation was made. More formally, SDMs can be interpreted as having the form $P(S|E)$, where S denotes the presence of a species, and E is an array of variables representing the local state of the environment at the point where the prediction is made.

As Franklin (2010a) highlights, SDMs emerged at a time where access to computers *and* the ability to effectively program them became easier. Although ecological insights, statistical methods, and data were pre-existing, the ability to turn these ingredients into something predictive required what is now called “computational literacy” – the ability to abstract, and automate, a system in order to generate predictions through computer simulations and their validation. One of the strengths of SDMs is that they can be used either for predictions or explanations (Elith & Leathwick 2009). To calculate $P(S|E)$ is to make a prediction (what are the chances of observing species S at a given location), that can be refined, validated, or rejected based on sampling; to understand what goes into E , *i.e.* what aspects of the environment are involved in determining the presence of a species, is an explanation of its distribution which can be related to, or informed by, knowledge from the species’ natural history.

Austin (2002) makes the point that SDMs exist at the interface between ecological theory and statistical models – this is a key feature of computational thinking, *i.e.* being able to integrate ideas and knowledge (abstract concepts) with statistical and numerical tools. In fact, one of the most recent and most stimulating developments in the field of SDMs is to refine their predictions not through the addition of more data, but through the addition of more processes (Franklin 2010b). These SDMs rely on the usual statistical models, but also on dynamical models (*i.e.* simulations; see *e.g.* Wisz et al. (2012) or Pellissier et al. (2013) for biotic interactions, and Miller & Holloway (2015) for movement and dispersal). What they lack in mathematical expressiveness (which is most often ruled out by the use of stochastic simulations), they assume to gain in predictive ability through the explicit consideration of more realistic ecological mechanisms (D’Amen et al. 2017; Staniczenko et al. 2017).

2 COMPUTATIONAL ECOLOGY IN ITS BROADER LANDSCAPE

2.1 *The four quadrats of ecological research*

In fig. 1, I propose a rough outline of four quadrats for ecological research. The first axis is based on the ability to *document* natural processes and their underlying mechanisms (through direct or indirect observation of natural systems) rather than *suggest* (through focus on a reduced number of mechanisms and their interactions). The second axis is based on the degree of integration between data and models, ranging from disconnected (for purely data-based or model-based) to highly integrated. A classification this coarse is bound to be caricatural, but it serves as an illustration of where computational ecology exists in the overall research methodology. Because it relies on the integration of data (if possible *raw* data from observational and manipulative experiments) and models (either statistical or phenomenological), it can *suggest* general trends through an abstraction of the idiosyncracies of a particular system.

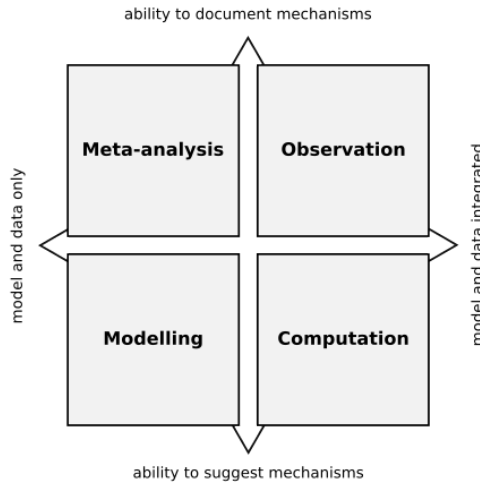


Figure 1 An overview of four quadrats of ecological research. The vertical axis differentiates the ability to document (by observation) or suggest (by simulation and inference) the action of ecological mechanisms. The horizontal axis indicates whether data and models are connected, or not. Computational ecology constitutes one of these quadrats, as it can bridge models with observations to further suggest mechanisms.

The specific example of predator-prey interactions should be a familiar illustration of how the same problem can be addressed in different ways. The classical prey-predator equations of Lotka & Volterra are an instance of a “modelling” based perspective, wherein mathematical analysis reveals how selected parameters (rates of interactions and growth) affect an ecologically relevant quantity (population stability and coexistence). These models, although they have been formulated to explain data generated through empirical observations, are disconnected from the data themselves. In fact, this family of model lies at the basis of a branch of ecological modelling that now exists entirely outside of data (Ackland & Gallagher 2004; Gyllenberg et al. 2006; Coville & Frederic 2013).

By contrast Sallan et al. (2011) study the same issue (sustained persistence and fluctuations of predator–prey couples through time) using a paleo-ecological timeseries, and interpret their data in the context of predictions from the Lotka-Volterra family of models (namely, they find support for Lotka-Volterra-like oscillations in time). Although dynamical models and empirical data interact in this example, they do not do so directly ; that is, the analysis of empirical data is done within the context of a broad family of model, but not coupled to *e.g.* additional simulations.

Meta-analyses, such as the one by Bolnick & Preisser (2005), are instead interested in collecting the outcome of observational and manipulative studies, and synthetizing the *effects* they report. These are often purely *statistical*, in that they aggregate significance, effect size, to measure how robust a result is across different systems. Meta-analyses most often require a *critical mass* of pre-existing papers (Lortie et al. 2013). Although they are irreplaceable as a tool to measure the strength of results, they are limited by their need for primary literature with experimental designs that are similar enough.

Where in this landscape does computational approaches fit?

2.2 Computational ecology in context

In *Life on the Mississippi*, Mark Twain wrote that “There is something fascinating about science. One gets such wholesale returns of conjecture out of such a trifling investment of fact”. This is a good description of the purpose of computational ecology: in a data-limited context, merging phenomenological models with pre-existing datasets is a way to efficiently develop conjectures, or more appropriately, build on our knowledge of models and data to put forward testable, quantified hypotheses. Pascual (2005) outlines that computational ecology has a unique ability to go from the complex (natural systems) to the simple (representations and conceptual models), and back (testable predictions). Although the natural world is immensely complex, it is paradoxically the high degree of abstraction of models that computational approaches favor that give them generality. Because the rate at which ecological data are collected is not improving, whereas our needs for testable and actionable predictions increases, refining the models and further integrating them with data is necessary.

This I illustrate in table 1. The quadrats of ecological approaches are ranked in (again, approximate and arbitrary) order of cost and effort. Ecological models make, by definition, high accuracy predictions, but they tend to be difficult to test (Rykiel 1996) – models relying on precise mathematical expressions can be difficult to calibrate or parameterize. Observations (field sampling) or manipulative approaches (micro/meso/macro-cosms, field experiments) are highly accurate (but have also immense human and monetary costs that limit the scale at which they can be applied). There is simply too much Nature around for us to observe, monitor, and manipulate it all.

Computational ecology ought to fall in the middle of this spectrum – although the reliance on advanced research computing incurs greater costs (either paid for by the researcher or by a computing center or consortium; these costs include operations of computing power, but also training of qualified students), it can generate predictions that are highly testable. Specifically, although the accuracy of these predictions is currently unknown (and will vary on a model/study/question basis), any additional empirical effort to *validate* predictions will improve

the predictions quality (Poisot et al. 2016).

Table 1 Overview of the properties of the quadrats delineated in fig. 1.

| Approach | accurate ? | testable ? | usable for prediction ? |
|-----------------------|------------|------------|-------------------------|
| Empirical observation | yes | | |
| Computational | unknown | yes | directly |
| Modelling | yes | no | indirectly |
| Meta-analysis | yes | no | no |

3 EN ROUTE TOWARDS SYNTHESIS

The field of ecology as a whole needs to improve the ways in which it can improve synthesis in order to become policy-relevant. Most of the global challenges have an ecological or environmental component, and outside of the socio-★ (ecological, economical, cultural, ...) aspects, ecologists can contribute to the mitigation or resolution of these challenges by i) assessing our knowledge of natural systems, ii) developing methods to produce scenarios using state-of-the-art models and tools, and iii) communicate the output of these scenarios to impact policy-making. White et al. (2015) propose that this falls under the umbrella of *action ecology*, i.e. using fundamental knowledge and ecological theory to address pressing, real-world questions.

Raghavan et al. (2016) suggest that this approach can also accommodate stakeholder knowledge and engagement. By building models that rely on ecological concepts, empirical data, and stakeholder feedback, they are able to implement a *computational agroecology* program, to use computational tools in the optimization of sustainable agricultural practices. This example suggests that not only can computational approaches yield fundamental research results in a short time frame, they can also be leveraged as a tool for applied research and knowledge transfer now. In the next sections, I will outline the key challenges that exist for the integration of computational ecology to other approaches to ecological research.

3.1 Mapping the domains of collaboration

Understanding how computational ecology will fit within the broader research landscape requires to answer three questions: what can computational ecology bring, what are the needs of computational ecologists, and what are the current limitations of computational approaches that could limit their immediate applicability. It seems, at this point, important to minimize neither the importance nor the efficiency of sampling and collection of additional data. Sampling is important because ecological questions, no matter how fundamental, ought to be grounded in phenomena happening in nature, and these are revealed by observation or manipulation of natural systems. Sampling is efficient because it is the final arbiter: how good any prediction is at explaining aspects of a particular empirical system is determined by observations of this system, compared to the predictions. Yet any endeavor to collect additional data has its scope limited by financial, human, and temporal constraints — or in other words, we need to chose what to sample, because we can't afford to sample it all. Computational approaches, because they can work through

large amounts of data, and integrate them with models that can generate predictions, might allow answering an all important question: what do we sample, and where? Some rely on their ecological intuition to answer; being a computational ecologist, and therefore being deprived of such intuitions, I would suggest that data, coupled with model, can be used to provide this answer.

Relying heavily on external information implies that computational research is dependant on standard for data representation. The Ecological Metadata Language (Fegeaus et al. 2005) is an attempt at standardizing the way meta-data are represented for ecological data; adherence to this standard, although it has been shown to improve the ease of assembling large datasets from single studies (Gil et al. 2011), is done on a voluntary basis (and is therefore abysmal). An alternative approach is to rely on community efforts to pre-curate and pre-catalog ecological data, such as with the flagship effort *EcoDataRetriever* (Morris & White 2013). Yet even this approach is ultimately limited, because of the human factor involved — when the upstream data change, they have to be re-worked into the software. A community consensus on data representation, although unlikely, would actually solve several problems at once. First, it would make the integration of multiple data sources trivial. Second, it will provide clear guidelines about the input and storage of data, thus maybe improving their currently limited longevity (Vines et al. 2014). Finally, it would facilitate the integration of data and models with minimum efforts and risk of mis-communication, since the format would be the same for all. To this extent, a recent proposal by Ovaskainen et al. (2017) is particularly interesting: rather than deciding on formats based on knowledge of eco-informatics or data management best practices, why not start from the ecological concepts, and translate them in digital representation? This task requires a strong collaboration between ecologists with topic expertise, ecologists with field expertise, and those of us leaning closest to the computational part of the field.

With or without a common data format, the problem remains that we have very limited insights into error propagation of predictions made on synthetic datasets (Poisot et al. 2016). There are biases in the underlying data, biases in the models used to generate the predictions, and this can turn out in three possible ways. First, predictions from these datasets accumulate bias and cannot be used. Second, because the scale at which these predictions are expressed is large, errors are (quantitatively) small enough to be over-ridden by the magnitude of actual variation. Finally, in the best-case but low-realism scenario, errors end up cancelling each other out. The best possible way to understand how errors propagate is to validate predictions *de novo*. Model-validation methods can be used, as they are with SDMS (Hijmans 2012), but *de novo* sampling carries the additional weight of being an independent attempt at testing the prediction. Improved collaborations on this aspect will provide estimates of the robustness of the predictions, in addition to highlighting the steps of the process in which uncertainty is high — these steps are natural candidates for additional methodological development.

Finally, there is a need to assess how the predictions made by purely computational approaches will be fed back into other types of research. This is notably true when presenting these approaches to stakeholders. One possible way to make this knowledge transfer process easier is to be transparent about the way predictions were derived: which data were used (with citations and unique identifiers), which software was used (with versions numbers and code), and what the model / simulations do. In short, the onus is on practitioners of computational research to make sure we provide all the information needed to communicate how predictions came to be.

3.2 *Establishing the currencies of collaboration*

An important question to further the integration of computational approaches to the workflow of ecological research is to establish *currencies* for collaborations. Both at the scale of individuals researchers, research groups, and larger research communities, it is important to understand what each can contribute to the research effort. As ecological research is expected to be increasingly predictive and policy-relevant, and as fundamental research tends to tackle increasingly refined and complex questions, it is expected that research problems will become more difficult to resolve; this is, I think, an incentive for a collaboration that builds on the skills that are specific to different approaches.

In an editorial to the *New England Journal of Medicine*, Longo & Drazen (2016) characterized scientists using previously published data as “research parasites” (backlash by a large part of the scientific community caused one of the authors to later retract the statement – Drazen (2016)). Although community ecologists would have, anyways, realized that the presence of parasites indicates a healthy ecosystem (Marcogliese 2005; Hudson et al. 2006), this feeling of unfair benefit to data re-analysis which is also expressed by ecologists (Mills et al. 2015) has to be addressed. It has no empirical support: Evans (2016) shows that the rate of data re-use in ecology is low and has a large delay – he found no instances of re-analysing the same data for the same (or similar) purpose. There is a necessary delay between the moment data are available, and the moment where they are re-used (especially considering that data are, at the earliest, published at the same time as the paper). This delay is introduced by the need to understand the data, see how they can be combined, develop a research hypothesis, etc..

On the other hand, there are multiple instances of combining multiple datasets collected at different scales, to address an entirely different question (see GBIF 2016–2016-10-10T11:13:13+02:00 for an excellent showcase) – it is more likely than data re-use is done with the intent of exploring different questions. It is also worth remembering that ecology (macroecology and biogeography in particular) already benefit immensely from data re-use – data collected by citizen scientists are used to generate estimates of biodiversity distribution, but also set and refine conservation targets (Devictor et al. 2010); an overwhelming majority of our knowledge of bird richness and distribution comes from the *eBird* project (Sullivan et al. 2009, 2014), which is essentially fed by the unpaid work of citizen scientists.

With this in mind, there is no tip-toeing around the fact that computational ecologists will be *data consumers*, and this data will have to come from ecologists with active field programs (in addition to government, industry, and citizens). Recognizing that computational ecology *needs* this data as a condition for its continued existence and relevance should motivate the establishment of a way to credit and recognize the role of *data producers* (which is discussed in Poisot et al. 2016, in particular in the context of massive dataset aggregation). Data re-users must be extremely pro-active in the establishment of crediting mechanisms for data producers; as the availability of these data is crucial to computational approaches, and as we do not share any of the cost of collecting these data, it behooves us to make sure that our research practices do not accrue a cost for our colleagues with field or lab programs.

3.3 *Training and advising computational ecologists*

The fact that data re-use is not instantaneous conveniently reveals another piece of information about computational ecology: it relies on different skills, and different tools. One of the most fruitful avenue for collaboration lies in recognizing the strengths of different domains: the skills required to assemble a dataset (taxonomic expertise, natural history knowledge, field know-how) and the skills required to develop robust computational studies (programming, applied mathematics) are different. If anything, this calls for increased collaboration, where these approaches are put to work in complementarity.

Barraquand et al. (2014) highlighted the fact that, professional ecologists received *less* quantitative and computational thinking that they think should be necessary. Increasing the amount of such training does not necessarily imply that natural history or field practice will be sacrificed on the altar of mathematics: rather, ecology would benefit from introducing more quantitative skills and reasoning across all courses, and introductory ones in particular (Hoffman et al. 2016). Instead of dividing the field further between empirically and theoretically minded scientists, this would showcase quantitative skills are being transversal to all questions that ecology can adress. What to teach, and how to integrate it to the existing curriculum, does of course requires discussion and consensus building by the community.

A related problem is that most practicing ecologists are terrible role models when it comes to showcasing good practices of data management (because there are no incentives to do this); and data management is a crucial step towards easier computational approaches. Even in the minority of cases where ecologists do share their data on public platforms, there are so few metadata that not being able to reproduce the original study is the rule (Roche et al. 2014, 2015). This is a worrying trend, because data management affects how easily research is done, regardless of whether the data are ultimately archived. Because the volume and variety of data we can collect tends to increase over time, and because we expect higher standard of analysis (therefore requiring more programmatic approaches), data management has already became a core skill for ecologists.

This view is echoed in recent proposals. Mislan et al. (2016) suggested that highlighting the importance of code in most ecological studies would be a way to bring the community to adopt higher standards, all the while de-mystifying the process of producing code. This also requires to equip ecologists with ways to evaluate the quality of the software they use (Poisot 2015). Finally, Hampton et al. (2015) proposed that the “Tao of Open Science” would be particularly beneficial to the entire field of ecology; as part of the important changes in attitude, they identified the solicitation and integration of productive feedback throughout the research process. Regardless of the technical solution, this emphasizes the need to foster, in ecologists in training, a culture of discussion across disciplinary boundaries.

4 CONCLUDING REMARKS

None of these approaches to ecological research have any intrinsic superiority – in the end, direct observation and experimentation trumps all, and serve as the validation, rejection, or refinement of predictions derived in other ways. The growing computational power, growing amount of

data, and increasing computational literacy means that producing theory and predictions is becoming cheaper and faster (regardless of the quality of these products). Yet the time needed to test any prediction is not decreasing, or not as fast. Computational science has resulted in the development of many tools and approaches that can be useful to ecology, since they allow to wade through these predictions and data. Confronting theoretical predictions to data is a requirement, if not the core, of ecological synthesis; this is only possible under the conditions that ecologists engage in meaningful dialogue, and recognize the currencies of their collaborations.

Finally, having this discussion about the place of computational ecology within the broader context of the ecological sciences will highlight areas of collaborations with other areas of science. Thessen (2016) makes the point that long-standing ecological problems would gain from being examined through a variety of techniques from the field of machine learning – I fully concur, because these techniques usually make the most of existing data (Halevy et al. 2009). Reaching a point where these methods are routinely used by ecologists will require a shift in our culture: quantitative training is perceived as inadequate (Barraquand et al. 2014), and most graduate programs do not train ecology students in contemporary statistics (Touchon & McCoy 2016).

Acknowledgements: I thank Dr. Allison Barner and Dr. Andrew McDonald for stimulating discussions, students of the Computational Ecology Summer School 2016 for asking “What is computational ecology?”, and the Station de Biologie des Laurentides de l’Université de Montréal for hosting me during part of the writing process. Part of this manuscript was inspired by presentations I gave during High Performance Computing Symposium 2015 in Montréal, for the organization of which I thank Compute Québec and Compute Canada, and during the 2016 meeting of the Ecological Society of America. I thank the volunteers of Software Carpentry and Data Carpentry, whose work contribute to improving the skills of ecologists.

REFERENCES

- Ackland & Gallagher.** (2004). Stabilization of Large Generalized Lotka-Volterra Foodwebs By Evolutionary Feedback. *Phys Rev Lett.* 93.
- Austin.** (2002). Spatial prediction of species distribution: an interface between ecological theory and statistical modelling. *Ecological Modelling.* 157:101–18.
- Barraquand et al.** (2014). Lack of quantitative training among early-career ecologists: a survey of the problem and potential solutions. *PeerJ.* 2:e285.
- Beaumont.** (2010). Approximate Bayesian Computation in Evolution and Ecology. *Annual Review of Ecology, Evolution, and Systematics.* 41:379–406.
- Bolker.** (2008). Ecological models and data in R. Princeton University Press;
- Bolnick & Preisser.** (2005). RESOURCE COMPETITION MODIFIES THE STRENGTH OF TRAIT-MEDIATED PREDATOR–PREY INTERACTIONS: A META-ANALYSIS. *Ecology.* 86:2771–9.
- Bourne.** (2011). Ten Simple Rules for Getting Ahead as a Computational Biologist in Academia.

PLoS Comput Biol. 7:e1002001.

Coville & Frederic. (2013). Convergence To The Equilibrium In A Lotka-Volterra Ode Competition System With Mutations. *arXiv:13016237*.

Devictor et al. (2010). Beyond scarcity: citizen science programmes as useful tools for conservation biogeography. *Diversity and distributions.* 16:354–62.

Drazen. (2016). Data Sharing and the Journal. *New England Journal of Medicine.* 374:e24.

D'Amen et al. (2017). Improving spatial predictions of taxonomic, functional and phylogenetic diversity. *Journal of Ecology.*

Elith & Leathwick. (2009). Species Distribution Models: Ecological Explanation and Prediction Across Space and Time. *Annu Rev Ecol Evol Syst.* 40:677–97.

Evans. (2016). Gauging the Purported Costs of Public Data Archiving for Long-Term Population Studies. *PLOS Biology.* 14:e1002432.

Fawcett & Higginson. (2012). Heavy use of equations impedes communication among biologists. *Proceedings of the National Academy of Sciences.* 109:11735–9.

Fegraus et al. (2005). Maximizing the Value of Ecological Data with Structured Metadata: An Introduction to Ecological Metadata Language (EML) and Principles for Metadata Creation. *Bulletin of the Ecological Society of America.* 86:158–68.

Franklin. (2010a). Mapping species distributions: spatial inference and prediction. Cambridge University Press;

Franklin. (2010b). Moving beyond static species distribution models in support of conservation biogeography. *Diversity and Distributions.* 16:321–30.

GBIF. 20162016-10-10T11:13:13+02:00. GBIF Science Review 2016 [Internet].

Gil et al. (2011). Examples of ecological data synthesis driven by rich metadata, and practical guidelines to use the Ecological Metadata Language specification to this end. *International Journal of Metadata, Semantics and Ontologies.* 6:46.

Gyllenberg et al. (2006). Limit cycles for competitor–competitor–mutualist Lotka–Volterra systems. *Physica D: Nonlinear Phenomena.* 221:135–45.

Halevy et al. (2009). The Unreasonable Effectiveness of Data. *IEEE Intelligent Systems.* 24:8–12.

Hampton et al. (2015). The Tao of open science for ecology. *Ecosphere.* 6:1–13.

Hijmans. (2012). Cross-validation of species distribution models: removing spatial sorting bias and calibration with a null model. *Ecology.* 93:679–88.

Hoffman et al. (2016). Development and Assessment of Modules to Integrate Quantitative Skills

in Introductory Biology Courses. *Cell Biology Education*. 15:ar14–4.

Hudson et al. (2006). Is a healthy ecosystem one that is rich in parasites? *Trends in ecology & evolution*. 21:381–5.

Legendre & Legendre. (1998). Numerical ecology. Oxford, UK: Elsevier;

Levin. (2012). Towards the marriage of theory and data. *Interface Focus*. 2:141–3.

Longo & Drazen. (2016). Data Sharing. *New England Journal of Medicine*. 374:276–7.

Lortie et al. 2013 Jun. Practical interpretation of ecological meta-analyses [Internet]. PeerJ PrePrints; Report No.: e38v1.

Marcogliese. (2005). Parasites of the superorganism: Are they indicators of ecosystem health? *International journal for parasitology*. 35:705–16.

Markowetz. (2017). All biology is computational biology. *PLOS Biology*. 15:e2002050.

May. (2004). Uses and Abuses of Mathematics in Biology. *Science*. 303:790–3.

Miller & Holloway. (2015). Incorporating movement in species distribution models. *Progress in Physical Geography*. 39:837–49.

Mills et al. (2015). Archiving Primary Data: Solutions for Long-Term Studies. *Trends in Ecology & Evolution*. 30:581–9.

Mislan et al. (2016). Elevating The Status of Code in Ecology. *Trends in Ecology & Evolution*. 31:4–7.

Morris & White. (2013). The EcoData Retriever: Improving Access to Existing Ecological Data. *PLoS ONE*. 8:e65848.

Otto & Day. (2007). A biologist's guide to mathematical modeling in ecology and evolution. Princeton University Press;

Ovaskainen et al. (2017). How to make more out of community data? A conceptual framework and its implementation as models and software. *Ecol Lett.*:n/a–a.

Papert. (1996). An exploration in the space of mathematics educations. *International Journal of Computers for Mathematical Learning*. 1.

Pascual. (2005). Computational Ecology: From the Complex to the Simple and Back. *PLoS Comp Biol*. 1:e18.

Pellissier et al. (2013). Combining food web and species distribution models for improved community projections. *Ecol Evol*. 3:4572–83.

Petrovskii & Petrovskaya. (2012). Computational ecology as an emerging science. *Interface*

Focus. 2:241–54.

Poisot. (2015). Best publishing practices to improve user confidence in scientific software. *Ideas in Ecology and Evolution*. 8.

Poisot et al. (2016). Synthetic datasets and community tools for the rapid testing of ecological hypotheses. *Ecography*. 39:402–8.

Raghavan et al. (2016). Computational Agroecology. *Proceedings of the 2016 CHI Conference Extended Abstracts on Human Factors in Computing Systems - CHI EA '16*. Association for Computing Machinery (ACM);

Roche et al. (2015). Public Data Archiving in Ecology and Evolution: How Well Are We Doing? *PLOS Biology*. 13:e1002295.

Roche et al. (2014). Troubleshooting Public Data Archiving: Suggestions to Increase Participation. Eisen, ed. *PLoS Biology*. 12:e1001779.

Rykiel. (1996). Testing ecological models: the meaning of validation. *Ecological Modelling*. 90:229–44.

Sallan et al. (2011). Persistent predator-prey dynamics revealed by mass extinction. *Proceedings of the National Academy of Sciences*. 108:8335–8.

Soetaert & Herman. (2008). A Practical Guide to Ecological Modelling: Using R as a Simulation Platform. Springer Verlag;

Staniczenko et al. (2017). Linking macroecology and community ecology: refining predictions of species distributions using biotic interaction networks. *Ecol Lett.*:n/a–a.

Sullivan et al. (2014). The eBird enterprise: an integrated approach to development and application of citizen science. *Biological Conservation*. 169:31–40.

Sullivan et al. (2009). eBird: A citizen-based bird observation network in the biological sciences. *Biological Conservation*. 142:2282–92.

Thessen. (2016). Adoption of Machine Learning Techniques in Ecology and Earth Science. *One Ecosystem*. 1:e8621.

Touchon & McCoy. (2016). The mismatch between current statistical practice and doctoral training in ecology. *Ecosphere*. 7:e01394.

Vines et al. (2014). The Availability of Research Data Declines Rapidly with Article Age. *Current Biology*. 24:94–7.

Warton et al. (2014). Model-based thinking for community ecology. *Plant Ecol*. 216:669–82.

White et al. (2015). The next generation of action ecology: novel approaches towards global ecological research. *Ecosphere*. 6:art134.

Wisz et al. (2012). The role of biotic interactions in shaping distributions and realised assemblages

of species: implications for species distribution modelling. *Biological Reviews*. 88:15–30.

Zhang. (2010). Computational ecology: artificial neural networks and their applications. Singapore: World Scientific Publ;

Zhang. (2012). Computational ecology: graphs, networks and agent-based modeling. New Jersey: World Scientific;